

Inference of Protein Secondary Structure with HMMs and BayesFlow

Simulation-Based Inference

Armin Maddah Asl Samaneh Naghdbishi Kiana Ebrahimi Farbod Geri Lila

July 16, 2026

TU Dortmund University

Course Instructor: Prof. Dr. Paul Bürkner

Introduction & Modeling Idea

The Problem

Predict per-position secondary structure from variable-length amino acid sequences.

Two-State HMM


State 0 = α -helix

State 1 = other (β -sheets & coils)

Modeling Goal

Approximate $P(S_t = \alpha \mid \mathbf{y})$ at every position t using **BayesFlow** neural posterior estimation.

Overall Pipeline




Validation

Compare to experimental ground truth:

Human insulin **1A7F** (chain A: 21 aa, chain B: 29 aa).

The 2-State Hidden Markov Model

State Transitions



Deterministic start: $P(S_1 = \text{other}) = 1.0$, $P(S_1 = \alpha) = 0.0$

Highly persistent: α stays α (90%), other stays other (95%)

Selected Emission Contrasts

Amino acid	α -helix	other	Role
A (Alanine)	12%	6%	Strong former
E (Glutamic)	9%	5%	Former
P (Proline)	2%	6%	Breaker
G (Glycine)	4%	9%	Breaker

Full 20×2 table in written report.

Prior Dynamics

Stationary distribution:

$$\pi(\alpha) = 0.05 / (0.05 + 0.10) = 1/3 \approx 33\%$$

$$\pi(\text{other}) \approx 67\%$$

Prior strongly favors **non-helical** states ($\sim 67\%$).

Data Generation & Analytical Targets

Simulator

Sequential Markov sampling of hidden states.

Independent randomized emission draws.

Random sequence lengths: 10–60 aa.

100% simulated — no real proteins.

Target Computation

Exact posteriors via Forward-Backward.

`hmmlearn.CategoricalHMM` — manually

configured, never `.fit()`.

Verification

Targets verified against independent

pure-NumPy FB solver. Confirms:

$P(S_0 = \alpha) \equiv 0$ (position 0 = seq. position 1).

Real-world validation: human insulin **1A7F** prepared here — sequence + experimental α -labels (used in slide 11).

On-demand Generation

`generate_dataset(n, ...)` produces

any number of (sequence, posterior,

states) tuples. Downstream code

chooses the split sizes.

Variable-Length by Design

Sequences kept **ragged** (10–60 aa),
each stored with its own length.

The data layer does **not pad** — padding
& masking are handled downstream by


the summary network.

Storage

`.npz` format: concatenated arrays +

`lengths` vector. No object arrays.


The Summary Network — BiLSTM



1. BiLSTM captures **sequential** dependencies — permutation-invariant nets fail here.
2. Padding token 20 & masked loss → handles variable-length sequences.
3. Two-stage: BiLSTM pretrained on BCE → frozen → CouplingFlow on top.

Two-Stage Training & CouplingFlow

Training Pipeline



Key Design Decisions

1. BiLSTM pretrained on masked BCE, then **frozen** → condition (dim 124 = 64 pooled + 60 logits).
2. CouplingFlow: depth 10, hidden 128, affine, actnorm → posterior draws.
3. **Model selection:** **not** flow val_loss → rank by masked MAE & correlation (selected depth 10).

Diagnostics: Adapted Simulation-Based Calibration

Why Standard SBC Does Not Apply

Standard SBC (Talts et al., 2018):

$\theta \sim p(\theta)$ (draw parameter)

$x \sim p(x | \theta)$ (simulate data)

Rank θ_{true} in posterior

→ Ranks must be **uniform**.

The Problem


HMM parameters are **fixed**.

$P(\alpha | \mathbf{y})$ is a **deterministic** function of the sequence.

No θ to sample from a prior.

→ Standard SBC is impossible.

Our Adaptation: Position-wise SBC





Key idea: Instead of ranking an unknown parameter, rank the **known** FB target within the flow's posterior draws at **each sequence position**.

6,764 positions tested → rank histogram, PIT ECDF, KS statistic.

Tests: Does the flow's posterior **envelope** the true value uniformly?

≡ Is the posterior **calibrated per quantile**?

Calibration Results — Adapted SBC Diagnostics



Rank histogram (left)

PIT ECDF with KS statistic (right)

6,764 positions


Adapted SBC Results

- Mean Rank = **0.65** (Rightward skew)
- Only **27%** of ranks below 0.50
- KS Distance = **0.234** (Highly significant)

The Spatial Cause

Position 0 (deterministic start) and padded boundaries are **exact zeros**. Normalizing flows are continuous — they struggle to represent **point masses**.

Calibration Results — Z-Score & Coverage Diagnostics



Key Metrics

Raw Z-std = **71.4**

(pos-0 degeneracy)

Trimmed (1.1%):

mean = **0.63**

std = **1.27**

Coverage:

87% vs. 90%



Overconfidence

Biased mean (0.63)

+ std > 1.0 (1.27)

→ posteriors **too narrow** at boundaries.

Synthetic Evaluation & Posterior Contraction



Black = True $P(\alpha)$ Blue band = BayesFlow 90% CI Red dashes = BiLSTM point

Accuracy Metrics

Model	MAE	r
Baseline (sequence-blind)	0.162	0.200
BayesFlow (median)	0.115	0.560
BiLSTM (point)	0.060	0.988

6,764 positions. Both beat the naive baseline.



Why BiLSTM Wins

$P(\alpha | \mathbf{y})$ is a **deterministic** function of the sequence. A direct network learns it with near-zero error. A flow models a residual density.

Why We Still Use BayesFlow

BiLSTM yields **no uncertainty**. BayesFlow provides full amortized posteriors $\rightarrow$ credible intervals, SBC, z-scores, and coverage diagnostics.

Real-World Validation — Human Insulin (1A7F)



The Statistical Mechanism

Chain B — Agreement ($r=0.86$)

34.5% true helix $\approx \pi=33\%$ prior. Prior and evidence agree $\rightarrow$ high accuracy.

Chain A — Conflict ($r=0.17$)

61.9% true helix $\gg 33\%$ prior. Severe prior-data conflict over only 21 residues.

Both beat sequence-blind baseline (MAE=0.162, $r=0.200$).

Key Proof — Model Limitation

Exact FB also fails on Chain A

Running the **exact analytical Forward-Backward** solver on Chain A achieves only $r \approx 0.40$.

$\rightarrow$ The failure is a **model limitation** of the 2-state HMM, **not** a neural network issue.

No amount of training can exceed the HMM's ceiling.

Limitations, Fixes & Take-Home Messages

Limitation & Fix: The HMM

Expand to **3 states** (α , β , coil).

Treat transition probabilities as **inferable** parameters θ under Beta priors.

→ BayesFlow would then recover a genuine posterior over θ .

Fix: The CouplingFlow

Pass targets through a **logit transformation**

before training → strictly bound all draws in $(0, 1)$.

Take-Home Messages (TL;DR)

1. BiLSTM summary nets capture sequential dependencies where permutation-invariant nets fail.
2. BayesFlow successfully amortizes posterior estimation, enabling robust uncertainty quantification.
3. Real-world validation demonstrates the necessity of prior typicality checking — identify prior-data conflicts.

Group Members

Armin Maddah Asl

armin.maddah@tu-dortmund.de

Samaneh Naghdbishi

samaneh.naghdbishi@tu-dortmund.de

Kiana Ebrahimi Farbod

kiana.ebrahimi@tu-dortmund.de

Geri Lila

geri.lila@tu-dortmund.de

Code:

github.com/armin2080/

[protein-secondary-structure-npe](#)